



Information Retrieval in High Dimensional Data

Lab #7, 28.06.2018

CVXOPT

Task 1. Machine Learning tasks are typically thought of optimization problems, e.g. minimizing an error function or maximizing a probability. Ideally, the optimization problem turns out to be convex, which implies that any local minimum is the global minimum of the formulation. In the following, it will be assumed that you have some basic knowledge about convex optimization. The intention of this task is to familiarize ourselves with CVXOPT, one of the most-widely used convex optimization toolboxes. Note: If CVXOPT does not accept your NumPy arrays, try casting them to `double`.

- Go to cvxopt.org and follow the installation instructions for your distribution. For conda, you need to run

```
conda install -c conda-forge cvxopt
```
- Skim through the **Examples** section on cvxopt.org to get an overview of the functionality of the different solvers of CVXOPT.
- Implement a function `minsq` which expects a NumPy array `A` of shape `(m,n)` and a NumPy array `y` of shape `(m,)` as its arguments and returns a NumPy array `x` of shape `(n,)` that solves the following problem.

$$\min_{\mathbf{x}} \|\mathbf{Ax} - \mathbf{y}\|.$$

Test your function by feeding it with appropriate inputs and comparing the results with the ones you get by using `np.linalg.pinv`. Experiment by adding white Gaussian noise to `y`.

Helpful Python/NumPy functions

<code>from cvxopt import matrix, solvers</code>	Basic CVXOPT functionality
<code>solvers.qp</code>	Quadratic Programming
<code>numpy.ndarray.astype</code>	Array casting (use <code>'double'</code>)